

Assignment 3

Fall 2017

CS834 Introduction to Information Retrieval

Dr. Michael Nelson

Orkun Krand

November 9, 2017

1 Question 6.1

1.1 Question

Using the Wikipedia collection provided at the book website, create a sample of stem clusters by the following process:

1. Index the collection without stemming.
2. Identify the first 1,000 words (in alphabetical order) in the index.
3. Create stem classes by stemming these 1,000 words and recording which words become the same stem.
4. Compute association measures (Dice's coefficient) between all pairs of stems in each stem class. Compute co-occurrence at the document level.
5. Create stem clusters by thresholding the association measure. All terms that are still connected to each other form the clusters.

Compare the stem clusters to the stem classes in terms of size and the quality (in your opinion) of the groupings.

1.2 Methodology

[CPickle](#) was utilized to speed processes up. I created the `stem.py` script to define the initial stem classes from the list of 1,000 words. This list was created using the inverted index I created for assignment 2. Roughly the first 15,000 items in list were/had digits so I used entries between 15,000 and 16,000. I used a threshold of 0.1 to clean up my result. Any stemmed words that exists in each class with a score lower than 0.1 were removed from their class.

1.3 Results

Initially converting each letter to lowercase resulted in more stems than I expected. The script returned 70 stem classes. aln: aln, alnes

alo: aloe, aloes

all: alling, all, alles, alle

alm: alm, alms

alkaloid: alkaloids, alkaloid

alight: alight, alighted

alkalin: alkaline, alkalinity

alexandrin: alexandrines, alexandrine

allori: allori, allory

alley: alleys, alley

alfieri: alfieri, alfieris

alien: alienated, alienate, aliens, alien, alienation, alienates, alienating

alpin: alpin, alpine, alpine

alexandr: alexandr, alexandre

alfr: alfre, alfred, alfr

alt: alte, altes, alt

allergi: allergies, allergy

allig: alligator, alligators

altdorf: altdorfer, altdorf
alp: alpes, alps, alpe, alp
almon: almoners, almoner
allemand: allemandes, allemande
allel: alleles, allele
allegori: allegories, allegory
allo: allo, allos
alphabet: alphabetical, alphabetically, alphabet, alphabetized, alphabetic, alphabets
altitud: altitude, altitudes
allyl: allylic, allyl
alli: allies, allis, alli, allying, allied, allie, ally
allud: allude, alluded, alludes
almont: almont, almonte
aleutian: aleutians, aleutian
altar: altars, altar
almshous: almshouse, almshouses
algorithm: algorithmic, algorithms, algorithm
alter: altered, altering, alterations, alteration, alters, alter
altern: alternator, alternating, alternators, alternated, alternately, alternate, alternatives, alternating, alternatively, alternative, alternations, alternates
almoravid: almoravid, almoravids
almond: almond, almonds
allow: allows, allowing, allow, allowed, allowe, allowances, allowable, allowance
algebra: algebraic, algebras, algebraically, algebra
alfi: alfie, alfy
algar: algars, algar
alexi: alexie, alexi, alexis
allstar: allstar, allstars
allotrop: allotrope, allotropes
allus: allusion, allusions
allur: alluring, allure
algarv: algarve, algarves
alphen: alphen, alphenal
alleg: allegedly, alleging, allegations, alleged, allegation, alleges, allege
alloc: allocates, allocation, allocate, allocations, allocated, allocating
alik: alikes, alike
alex: alex, alexs
alma: alma, almas
alin: aline, alin
alkyl: alkyl, alkylating
alia: alia, alias
align: aligning, align, alignable, aligns, aligned, alignments, alignment
allround: allround, allrounder
alon: alone, alones, alon
alor: alor, alors
alias: aliasing, aliases
alouett: alouettes, alouette
allot: allotted, allot, allotment
almanac: almanacs, almanac
alloy: alloy, alloys, alloyed
alland: alland, allande

alga: alga, algae
allianc: alliance, alliances
allevi: alleviating, alleviate

2 Question 6.2

2.1 Question

Create a simple spelling corrector based on the noisy channel model. Use a single-word language model, and an error model where all errors with the same edit distance have the same probability. Only consider edit distances of 1 or 2. Implement your own edit distance calculator (example code can easily be found on the Web)

2.2 Methodology

The spelling corrector will be based on [Peter Norvig's corrector](#) he wrote while traveling on a plane. For the sample text I used [big.txt](#) which I found on Peter Norvig's website. The probability of a word is calculated with the following formula:

$$P(w) = \frac{Count_w}{N}$$

where $Count_w$ is the word count for word w and N is the sum of all word counts.

In order to determine if a word is right or wrong, I found all the words with edit distance one and two to each word. Assuming the correct word will be the one with shorter edit distance, those words with a shorter edit distance and higher probability are picked as the correct word.

The code can be found in `spell.py`.

2.3 Results

Here is a screenshot of a couple words I tested. If it can't come up with a replacement "correct" word, the script returns the original word.

```
Krando67:A3 Krando67$ python spell.py elizabet
elizabeth
Krando67:A3 Krando67$ python spell.py elizab
eliza
Krando67:A3 Krando67$ python spell.py eht
eh
Krando67:A3 Krando67$ python spell.py teh
the
Krando67:A3 Krando67$ python spell.py embrla
emboli
Krando67:A3 Krando67$ python spell.py embrlla
umbrella
```

Figure 1: Sample Spelling Correcter Output

3 Question 6.5

3.1 Question

Describe the snippet generation algorithm in Galago. Would this algorithm work well for pages with little text content? Describe in detail how you would modify the algorithm to improve it.

3.2 Answer

The main idea of the snippet generator is that it looks for the query words in the text of the document. Once a match is found, the generator grabs the 5 words before and after to create a snippet region. It can handle overlaps in snippet regions. This search goes on until the character threshold is reached at which point the generator stops going through the document. Snippet's ready.

This algorithm would fail to create a useful snippet for a page with little text content since it relies on document text to generate the snippet regions. If there are no words in the document, then the generator has nothing to compare the query words to.

Taking into account the query word density is probably a good idea as it would improve the results shown. So for example if a user is searching for the phrase "tropical fish", and the snippet generator finds the word tropical over and over again, that would be enough to generate a snippet. This will force the algorithm to run for longer, increasing overhead but maybe we would be using Galago today instead of [Google](#) if they cared to implement it. In order to implement this, snippet regions could be entered into a hash table which keeps track of which query words are in which snippet regions. Afterwards, the hash table would be analyzed to find which snippet regions have the highest density of query words. Those snippet regions would be concatenated to create the final snippet.

4 Question 6.9

4.1 Question

Give five examples of web page translation that you think is poor. Why do you think the translation failed?

4.2 Methodology

I will use the [Translate extension of Apple's Safari](#) (I don't normally use Safari but Chrome knows me too well) which uses Google or Bing to translate 3 pages from Turkish to English and two pages from English to Turkish.

4.3 Examples

4.3.1 [Motosiklet satışları düşüyor](#)

This is a news article about the low number of motorcycle sales in Turkey. While the translation is readable, it is very hard to understand. Some words are not translated at all. The first paragraph was going well until the last sentence where the translator butchered it quite bad. This was caused

by a word that has more than one meaning translated as one meaning when it should be translated as the other meaning. In this specific example, the word “adet” means each or copies. “Average 140 thousand copies” makes no sense since motorcycles aren’t made out of paper. Omitting the culprit would’ve probably solved the problem since in English, we don’t have to say “140 thousand each motorcycle”.

Also it translates the currency from Turkish Lira to British Pounds without changing the number. Considering the value difference between the two currencies, someone reading this will think Turkish people pay \$4,500 per year for motorcycle insurance when the actual number is \$900.

4.3.2 DÄŖZELLER DÄŖZELÄŖ DÄŖZELERÄŖ DÄŖZDÄŖ

“Sahibinden” is a platform for people to sell vehicles, houses, find jobs and so forth. It’s very similar to craigslist but it has a better interface. This ad is for a Opel Astra 1.3 diesel car. I will include screenshots in case the ad disappears.

Looking at the information about the vehicle, you will notice that some things don’t make any sense, like the color of the vehicle for example. The word they use in the original page “fÄŖijme” translates to smoked but is usually referred to the color you see in the picture. After checking out various dictionaries, I have no idea how they came up with the word “rendering”.

The best part of this ad is the description. The person who posted this ad decided to write a poem that describes his/her vehicle. I will let you try to make some sense of it. The main reason translation fails so bad is because the poem isn’t really written in a traditional sentence structure. It puts words in different orders to ensure it rhymes but this breaks the translator as I assume it looks at the previous and next word (or the whole sentence) when picking a translation for each word.


4.3.3 Lamborghini’den bÄŖijtÄŖge dostu araÄŖ mÄŖs geliyor?

This is a news article from a mainstream newspaper website (the first one was a motorcycle magazine’s website). This review is mostly bad because the algorithm is not good enough. It is a news article so it’s properly written (no spelling mistakes, well structured sentences). The translator just fails to recognize some words such as the word “demek” in the first sentence. It means to say something in English. It can also be thought of as calling something. So if we were to put the word at the very beginning of the sentence, that would make the first sentence perfect. It’s mostly small things like this that ruin the translation on this page.

4.3.4 Here Are Some Tips for Parting With a Car You Love


Similar to the problems with translating from Turkish to English, Bing’s English to Turkish translator has issues. The first sentence of the second paragraph doesn’t make much sense. We can guess that he probably owned it for 2 years. But the bigram “daily driver” is a unique because the word “driver” generally means someone that drives a vehicle. The translator seems to have an easier time with every day words compared to words like “bolstered” which nobody uses daily. When a sentence has more common words, it seems to be able to process more words at a time (bigrams for uncommon words vs. five-gram for more common words).

[Anasayfa](#) [Vasıta](#) [Otomobil](#) [Opel](#) [Astra](#) [1.3 CDTI](#) [Sport](#)

 **Bu ilan o kadar çok sevildi ki, Efsane İlan oldu**
Tüm Efsane İlanları Gör

DİZELLER DİZELİ DİZELERİ DİZDİ ★ Favorilerime Ekle Yazdır

#359464489





55.000 TL
Ankara / Çankaya / Çayyolu Mah.

İlan No	359464489
İlan Tarihi	01 Kasım 2016
Marka	Opel
Seri	Astra
Model	1.3 CDTI Sport
Yıl	2012
Yakıt	Dizel
Vites	Manuel
KM	65.000
Kasa Tipi	Sedan
Motor Gücü	95 hp
Motor Hacmi	1248 cc
Çekiş	Önden Çekiş
Renk	Füme
Garanti	Hayır
Plaka / Uyruk	Türkiye (TR) Plakalı
Kimden	Sahibinden
Takas	Hayır
Durumu	İkinci El

Güvenlik İpuçları
İlgilendiğiniz aracı görmeden ödeme yapmayın, para göndermeyin.
Detaylı bilgi için tıklayın.

BÖYLE BİR ARACI 899 TL'den BAŞLAYAN TAKSİTLE SAHİP OLABİLİRSİNİZ
HEMEN KEŞFET
fuzuloto

Büyük Fotoğraf Video



1/2 Fotoğraf

İlan Detayları **Teknik Özellikler**

Açıklama

İlan ile İlgili Şikayetim Var

Figure 2: Opel Astra Turkish

4.3.5 Norfolk, Virginia

I thought the translation of Wikipedia pages would be the worst which is why I saved it for last but this is a surprisingly good translation. I was able to observe how well the translator works from the previous examples but this is better compared to others. Maybe Google's better at translating from English than translating to English. But the main reason it's better in my opinion is because the Wikipedia page contains a lot of special names such as "Chesapeake", or "Norfolk Southern Railway". These help the reader understand the context even if other words fail to translate correctly.

Google Translated to: English Show original Options


sahibinden.com Kelime, ilan no veya mağaza adı ile ara Detailed Search Sign in Sign up

Home page Vehicle Car Opel Astra 1.3 CDTI Sport

This ad was so much loved that legend announcement was all the legend Ad Rates

DIESEL DIESEL strings were KNEE

#359464489



[Larger Photo](#) [Video](#)

1 / 2 Photos

55.000 TL

Ankara / Çankaya / Çayyolu Mah.

Listing ID 359464489

Ad Date November 1, 2016

Marka Opel

Seri Astra

Model 1.3 CDTI Sport

year 2012

Fuel Diesel

Vites Manuel

KM 65.000

Less tipi sedan

Engine power 95 hp

Engine Displacement 1248 cc

Traction Front Wheel Drive

Color Rendering

Warranty No

Sticker / Nationality Turkey (TR) Plate

From the owner

Clearing No.

Condition Second Hand

[There Ad-Related Complaints](#)

Ad Details **Technical Specifications**

Explanation

Safety Tips
without seeing the vehicle you making a payment as a depos money.
For more information, click here.

BÖYLE BİR ARA
899 TL'den
BAŞLAYAN TAKSİT
SAHİP OLABİLİR

HEMEN KEŞFE

fuzulot
KULAY OTOMOBİL EKONOMİK MİRAL

Compare

Figure 3: Opel Astra English

4.4 Discussion

The first two pages were translated from Turkish to English using Google Translate. The next two were translated using Bing and the last one was translated using Google as well. The translator seems to get confused by words that have more than one meanings. Poems definitely throw it off and provide a lower accuracy translation. Sentences that give data are better translated. The ones with numerical

ERROR NO MEANS Me, I CALLED NO FAULTS,
A SLICE OF LIES DÖNMEZ, I can NEVER SAID,
UNLAWFUL TOUCH is, I vow in the EEA can not EAT,
FIELD HAYIRLI, WHETHER UĞURLU.

PAINT is changing THERE? HAS.
I CARPOOL the MUL, REASON FOR CARD,
I TROUBLE USING THREE YEARS,
FIELD HAYIRLI, WHETHER UĞURLU.

CATS LEGEND VEHICLE HAD DEM was AFRAID,
It IS NOT SPEED, THIRTY MAYBE, MAYBE I KIRK,
He UNPAINTED THIS VEHICLE ACCIDENT FOR MY FIRST ONE was,
FIELD HAYIRLI, WHETHER UĞURLU.

DAMAGE THAT WOULD BE THE SECOND GEAR in HOW MUCH?
DAMAGE AND MAINTENANCE In AUTHORIZED SERVICE,
If FULL LIST WHAT MADE ALL REGISTERED,
FIELD HAYIRLI, WHETHER UĞURLU.

I do not YOGURT ANYONE SOUR DE is TRUE,
Whose of HEALING, some that are STOMACH PAIN,
ERROR EXPENSES, BUT EVEN CURVE FRAME,
FIELD HAYIRLI, WHETHER UĞURLU.

Figure 4: Car Ad Poem English

data are even better since the languages have pretty well defined rules as to what comes before or after numbers.

ARACIMDA HATA YOK, KUSUR YOK DİYEMEM,
YALANA DİLİM DÖNMEZ, ASLA SÖYLEYEMEM,
HARAM DOKUNUR, AÇ KALSAM DA YİYEMEM,
ALANA HAYIRLI, UĞURLU OLSUN.

BOYASI, DEĞİŞENİ VAR MIDIR ? VARDIR.
ÖNDEKİNE ÇARPMIŞTIM, SEBEBİ KARDIR,
SORUNSUZ KULLANIYORUM ÜÇ YILDIR,
ALANA HAYIRLI, UĞURLU OLSUN.

ARAÇTA EFSANE KEDİM DE VARDI, KORKTU,
HIZIM ÇOK DEĞİL, BELKİ OTUZ, BELKİ KIRKTI,
BOYASIZDI, ARACIM İÇİN BU KAZA BİR İLKTİ,
ALANA HAYIRLI, UĞURLU OLSUN.

NE KADAR HASAR OLUR Kİ İKİNCİ VİTESTE ?
HASARI VE BAKIMLARI YETKİLİ SERVİSTE,
NE YAPILDIYSA HEPSİ KAYITLI TAM LİSTE,
ALANA HAYIRLI, UĞURLU OLSUN.

KİMSE YOĞURDUM EKŞİ DEMEZ, DOĞRUDUR,
KİMİNİNKİ ŞİFA, KİMİNİNKİ MİDEDE AĞRIDIR,
HATASIZ DERLER, AMA ŞASESİ BİLE EĞRİDİR,
ALANA HAYIRLI, UĞURLU OLSUN.

Figure 5: Car Ad Poem Turkish

5 Question MLN2

5.1 Question

Using the small wikipedia example, choose 10 words and compute MIM, EMIM, chi square, dice association measures for full document & 5 word windows (cf. pp. 203-205)

5.2 Methodology

I created `mln2.py` which tackles this problem. It uses the same inverted index as 6.1 to calculate the measures requested. The words I picked were:

- umbrella
- messenger
- cramps
- equestrian
- sea
- everlasting
- association
- vehicle
- python
- motorcycle

5.3 Results

The words “python” and “everlasting” returned the same results for each type of association. X2 and Dice seem to be better compared to MIM and EMIM as they return related words. Dice performed better than X2 especially for the word “sea”. For that word, Dice was spot on, suggesting very related terms like island, ocean, ships, etc. Worst seems to be EMIM as it kept suggesting the same 8 words for the words “association”, “vehicle”, and “motorcycle”. None of those 8 words were related to the query word.

umbrella

MIM ['tumen', 'haoma', 'hyperdrive', 'spaceships', '1521ad', 'tuuli', 'b0007bfxk4', 'vedism', 'vorgejz', 'temp07']

EMIM ['which', '1', '2008', '4', 'on', 'categories', 'offline', 'gnu', '7ecommon', '501']

X2 ['tumen', 'haoma', 'hyperdrive', 'spaceships', '1521ad', 'tuuli', 'b0007bfxk4', 'vedism', 'vorgejz', 'temp07']

Dice ['confucianism', 'amendment', 'token', 'undesirable', 'disagreed', 'hostages', 'salvadoran', 'solemn', 'repel', 'unconstitutional']

messenger

MIM ['kouga', 'wftv1', 'susans', 'wflc', 'furry', 'wkhk', 'casbah', 'brookmeade', 'orseno', 'wedekind']

EMIM ['2008', '2', 'on', 'categories', 'offline', 'gnu', '7ecommon', '501', 's', 'fixalpha']

X2 ['salama', 'wcau', 'kouga', 'wftv1', 'susans', 'wflc', 'furry', 'wkhk', 'casbah', 'brookmeade']

Dice ['salama', 'wcau', 'copious', 'striver', 'personality', 'cared', 'jarvis', 'sanford', 'cochran', 'prophet']

cramps

MIM ['gastroenterology', 'boulardii', 'hydrophila', 'frailty', 'challis', 'earshot', 'mismanaged', 'anti-motility', 'paregoric', 'mongering']

EMIM ['unhygienic', 'buying', 'drinking', 'eating', 'bed', 'sullivan', 'blood', 'selling', 'getting', 'scale']

X2 ['unhygienic', 'gastroenterology', 'boulardii', 'hydrophila', 'frailty', 'challis', 'earshot', 'mismanaged', 'antimotility', 'paregoric']

Dice ['unhygienic', 'gastroenterology', 'boulardii', 'hydrophila', 'frailty', 'challis', 'earshot', 'mismanaged', 'antimotility', 'paregoric']

equestrian

MIM ['deligia', 'pimplo', 'irredento', 'liamodwyer13', 'louisay', 'corporacy', 'damiano', 'ongarato', 'morbidegli', 'maccabean']

EMIM ['olympics', 'summer', '2008', 'bronze', 'on', 'categories', 'offline', 'gnu', '7ecommon', '501']

X2 ['equestrianism', 'equestrians', 'deligia', 'pimplo', 'irredento', 'liamodwyer13', 'louisay', 'corporacy', 'damiano', 'ongarato']

Dice ['equestrianism', 'equestrians', 'archery', 'fencing', 'deligia', 'pimplo', 'irredento', 'liamodwyer13', 'louisay', 'corporacy']

sea

MIM ['kalecik', 'vologda', 'immunities', 'rbeas', 'hoogenboom', 'antillaise', 'krais', 'wheatear', 'balade']

EMIM ['categories', 'offline', 'gnu', '7ecommon', '501', 's', 'fixalpha', '7emonobook', 'cdata', 'disclaimers']

X2 ['2', '4', '5', 'a', 'showtoctoggle', 'tocshowtext', 'tochidetext', '8', 'show', '6']

Dice ['ocean', 'h', 'over', 'island', 'earlier', 'little', 'atlantic', 'ship', 'bay', 'islands']

everlasting

MIM ['andy120', 'didactohedron', 'prostration', 'auricular', 'oris', 'admonish', 'starets', 'martyrion', 'deinde', 'divulges']

EMIM ['andy120', 'didactohedron', 'prostration', 'auricular', 'oris', 'admonish', 'starets', 'martyrion', 'deinde', 'divulges']

X2 ['andy120', 'didactohedron', 'prostration', 'auricular', 'oris', 'admonish', 'starets', 'martyrion', 'deinde', 'divulges']

Dice ['andy120', 'didactohedron', 'prostration', 'auricular', 'oris', 'admonish', 'starets', 'martyrion', 'deinde', 'divulges']

association

MIM ['pskalka', 'foundation', 'vann', 'kovaciny', 'tcbr', 'immunities', 'landrover4', 'hoogenboom', 'sunbeams', 'kids']

EMIM ['offline', 'gnu', '7ecommon', '501', 's', 'fixalpha', '7emonobook', 'cdata', 'disclaimers', 'mediawiki']

X2 ['national', 'show', 'showtoctoggle', 'tocshowtext', 'tochidetext', '4', '1', '2', '5', '6']

Dice ['national', 'school', 'year', 'retrieved', '2002', '2000', 'team', 'non', 'showtoctoggle', 'tocshowtext']

vehicle

MIM ['rbeas', 'numerao', 'skatelikekemp', 'mklimstra', 'skyblue27', 'nichiyoubi', 'samokhodno', 'syx', 'kcci', 'conductive']

EMIM ['on', 'offline', 'gnu', '7ecommon', '501', 's', 'fixalpha', '7emonobook', 'cdata', 'disclaimers']

X2 ['vehicles', 'wheels', 'axle', 'motor', 'cars', 'engine', 'policemen', 'sfoskett', 'tracked', 'tanks']

Dice ['vehicles', 'motor', 'cars', 'wheels', 'engine', 'driving', 'mounted', 'car', 'fitted', 'electronic']

python

MIM ['kudrimalai', 'drypetes', 'elephas', 'trijuga', 'disopyros', 'manilkara', 'modesta', 'wcmc', 'shrew', 'dendrocygna']

EMIM ['kudrimalai', 'drypetes', 'elephas', 'trijuga', 'disopyros', 'manilkara', 'modesta', 'wcmc', 'shrew', 'dendrocygna']

X2 ['kudrimalai', 'drypetes', 'elephas', 'trijuga', 'disopyros', 'manilkara', 'modesta', 'wcmc', 'shrew', 'dendrocygna']

Dice ['kudrimalai', 'drypetes', 'elephas', 'trijuga', 'disopyros', 'manilkara', 'modesta', 'wcmc', 'shrew',

'dendrocygna']

motorcycle

MIM ['dnq', 'smackdown', 'jugulator', 'indyahh', 'seasalt', 'y2kcrazyjoker', 'derbi', 'toshihisa', 'downey-ocean', 'puniet']

EMIM ['categories', 'offline', 'gnu', '7ecommon', '501', 's', 'fixalpha', '7emonobook', 'cdata', 'disclaimers']

X2 ['motogp', 'puniet', 'guintoli', 'capirossi', 'gibernau', '125cc', 'loris', 'dovizioso', 'melandri']

Dice ['motogp', 'brianhe', 'fabrizio', 'cc', 'puniet', 'guintoli', 'capirossi', 'gibernau', '125cc', 'loris']